



Fuzzing

Yajin Zhou (<http://yajin.org>)

Zhejiang University



Testing

Testing is the process of executing a program to find errors.

An error is a deviation between observed behavior and specified behavior, i.e., a violation of the underlying specification:

- Functional requirements (features)
- Operational requirements (performance, usability)
- Security requirements



Testing

- Manual testing
 - Unit testing (individual modules)
 - Integration testing (interaction between modules)
 - System testing (full application testing)
- Fuzzing testing
- Symbolic and concolic testing

Fuzzing

- As Its Core, Fuzzing is Random Testing
 - it starts a long time ago
 - automated software testing technique

1981 Random testing is a cost-effective alternative to systematic testing techniques (Duran & Natos)

1983 "The Monkey" (Capps)

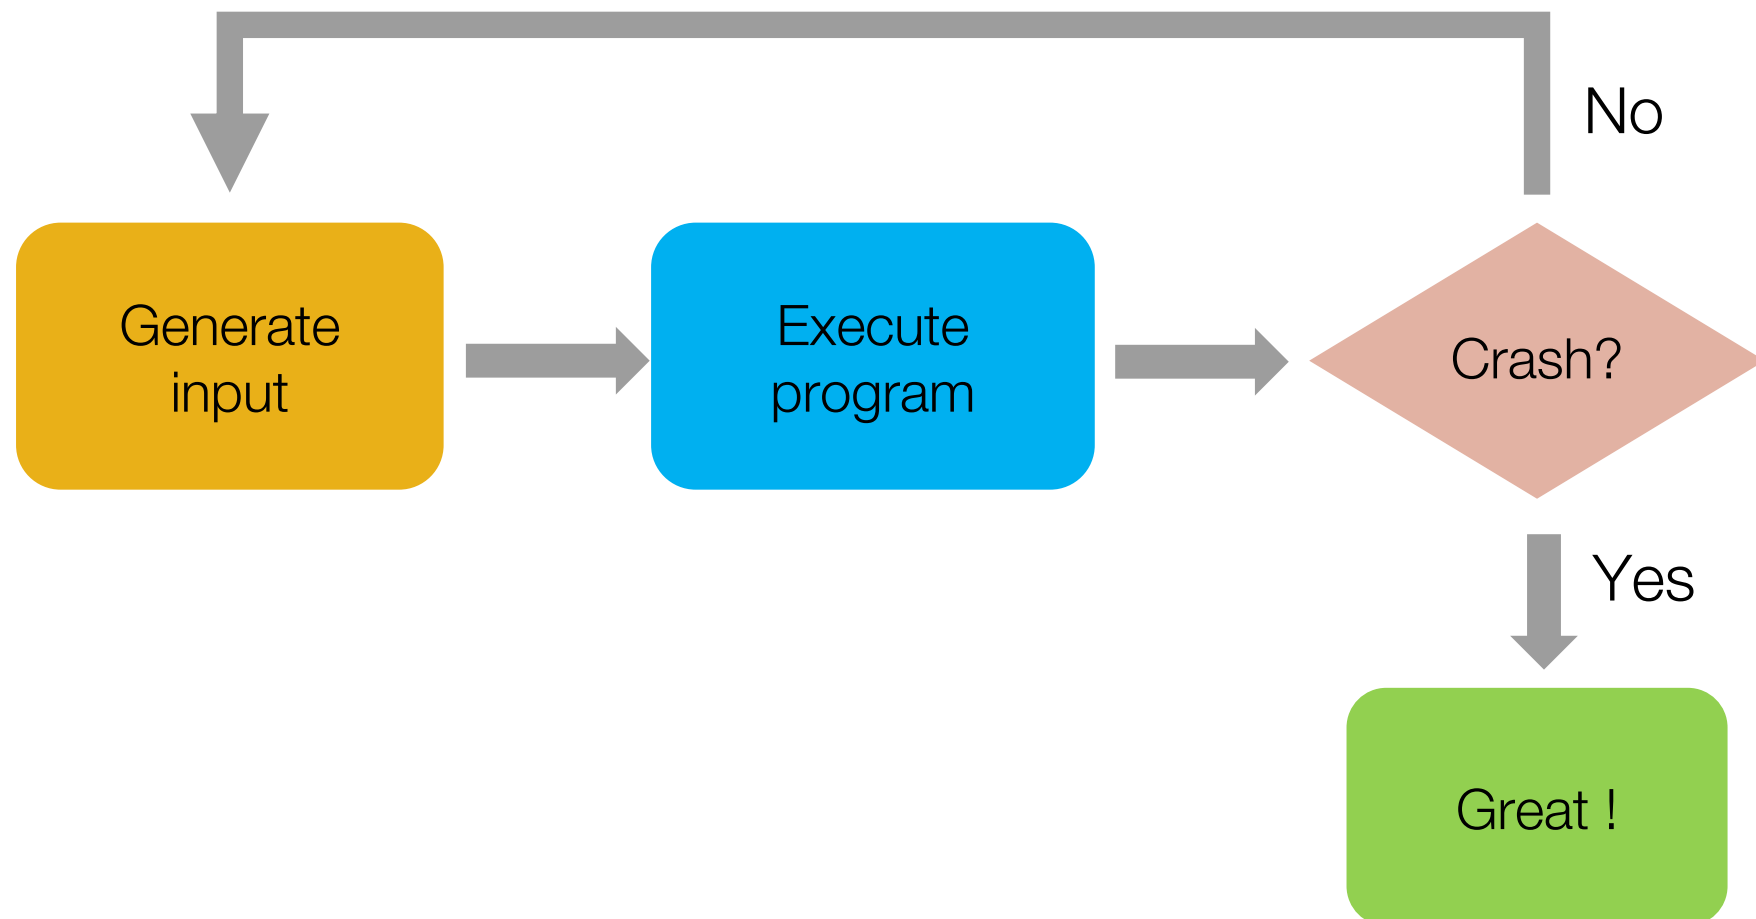
1988 Birth of the term "fuzzing" (Miller)



Monkey testing

Fuzzing

- Feeding in random inputs until the program crashes





Fuzzing

- Feeding in random inputs until the program crashes

- How do we inject inputs?
- How do we generate inputs?
- How do we automate the process?
- How do we execute the program?
- How do we detect bugs?



Fuzzing

- Feeding in random inputs until the XML parser crashes

- How do we inject inputs?
- How do we generate inputs?
- How do we detect bugs?
- How do we automate the process?



Fuzzing

- Feeding in random inputs until the XML parser crashes

— How do we inject inputs? Execute the parser with a xml file

— How do we generate inputs?

— How do we detect bugs?

— How do we automate the process?



Fuzzing

- Feeding in random inputs until the XML parser crashes

— How do we inject inputs? execute the parser with a xml file

— How do we generate inputs?

— How do we detect bugs?

— How do we automate the process?



Generating inputs for programs

- In case of an XML parser

Idea #1: just generate random binary data

```
cat /dev/urandom | xml_parser
```



Random inputs

```
if (input[0] == '<')  
  
    if (input[1] == 'x')  
  
        if (input[2] == 'm')  
  
            if (input[3] == 'l')  
  
                // start process file
```

- Parser expects the file to start with `<xml` header
- We need $\sim 2^{\{8\}^4}$ guesses to get past the header check
- works poorly & incomplete



Generating better inputs for programs

Idea #2: Model what the application should process


Structured inputs (a.k.a. structure-aware fuzzing)



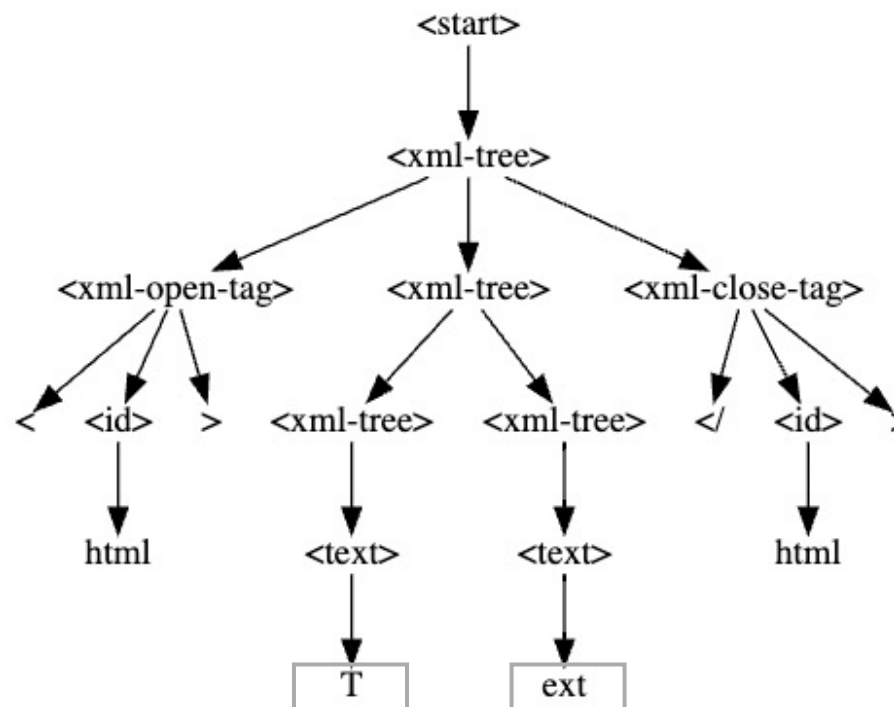
Structured inputs

```
XML_GRAMMAR: Grammar = {
    "<start>": ["<xml-tree>"],
    "<xml-tree>": ["<text>",
        "<xml-open-tag><xml-tree><xml-close-tag>",
        "<xml-openclose-tag>",
        "<xml-tree><xml-tree>"],
    "<xml-open-tag>": ["<<id>>", "<<id> <xml-attribute>>"],
    "<xml-openclose-tag>": ["<<id>/>", "<<id> <xml-attribute>/>"],
    "<xml-close-tag>": ["</<id>>"],
    "<xml-attribute>": ["<id>=<id>", "<xml-attribute> <xml-attribute>"],
    "<id>": ["<letter>", "<id><letter>"],
    "<text>": ["<text><letter_space>", "<letter_space>"],
    "<letter>": srange(string.ascii_letters + string.digits +
        "\"'."),
    "<letter_space>": srange(string.ascii_letters + string.digits +
        "\"' " + "\t"),
}
```

Structured inputs



Initial structured tree




structured tree after adding nodes



Structured inputs

After several
mutation



new structured tree



```
<start>
|-<html><head><title>Hello</title></head><body>World<br/></body></html>
<xml-tree>
|-<html><head><title>Hello</title></head><body>World<br/></body></html>
|-<head><title>Hello</title></head><body>World<br/></body>
|-<head><title>Hello</title></head>
|-<title>Hello</title>
|-Hello
|-<body>World<br/></body>
|-World<br/>
|-World
|-<br/>
<xml-open-tag>
|-<html>
|-<head>
|-<title>
|-<body>
<xml-openclose-tag>
|-<br/>
<xml-close-tag>
|-</title>
|-</head>
|-</body>
|-</html>
<xml-attribute>
<id>
<text>
<letter>
<letter_space>
```

generated testcase



Generating better inputs for programs

Idea #3: Coverage as completeness metric

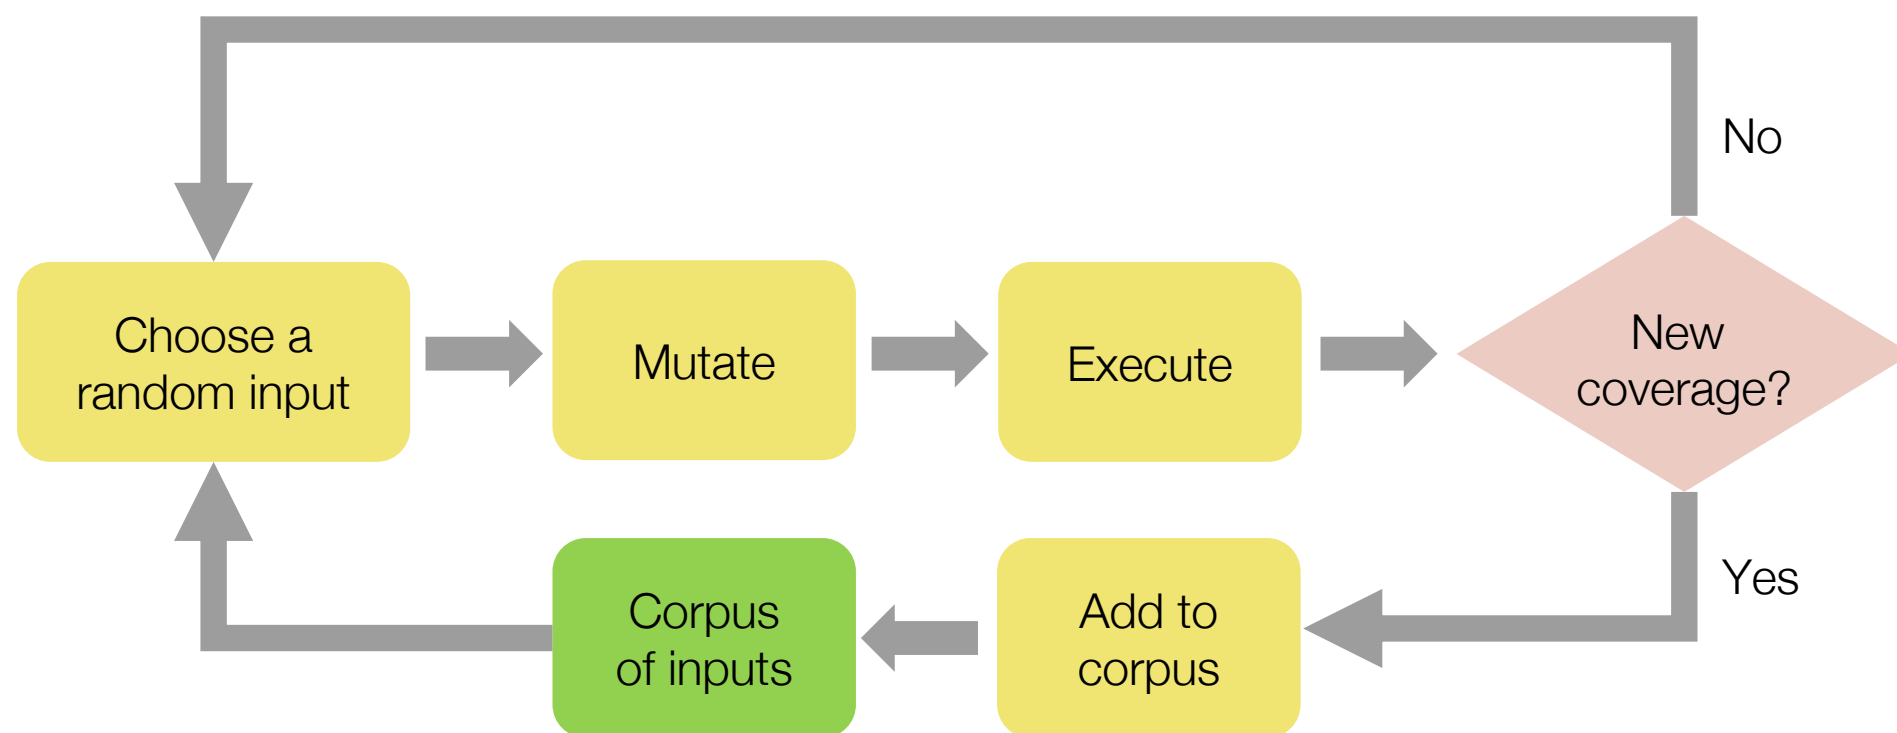
Intuition:

- A software flaw is only detected if the flawed statement is executed.
- Effectiveness of test suite therefore depends on how many statements are executed.

Coverage-guided generation (a.k.a. coverage-guided fuzzing)



Coverage-guided generation





Branch Coverage

```
int arr[6];
```

```
int func(int a, int b) {  
    int idx = 6;  
    if (a < 6) idx -= 6; else idx -= 1;  
    if (b < 6) idx -= 1; else idx += 1;  
    return arr[idx]; // idx = 4  
}
```

Test input: a=10, b=1



Branch Coverage

```
int arr[6];
```

```
int func(int a, int b) {  
    int idx = 6;  
    if (a < 6) idx -= 6; else idx -= 1;  
    if (b < 6) idx -= 1; else idx += 1;  
    return arr[idx]; // idx = 1  
}
```

Test input: a=1, b=10



Branch Coverage

```
int arr[6];
```

```
int func(int a, int b) {  
    int idx = 6;  
    if (a < 6) idx -= 6; else idx -= 1;  
    if (b < 6) idx -= 1; else idx += 1;  
    return arr[idx];  
}
```

All test inputs: $a=10, b=1$ and $a=1, b=10$

➡ Full branch coverage



Is branch coverage enough?

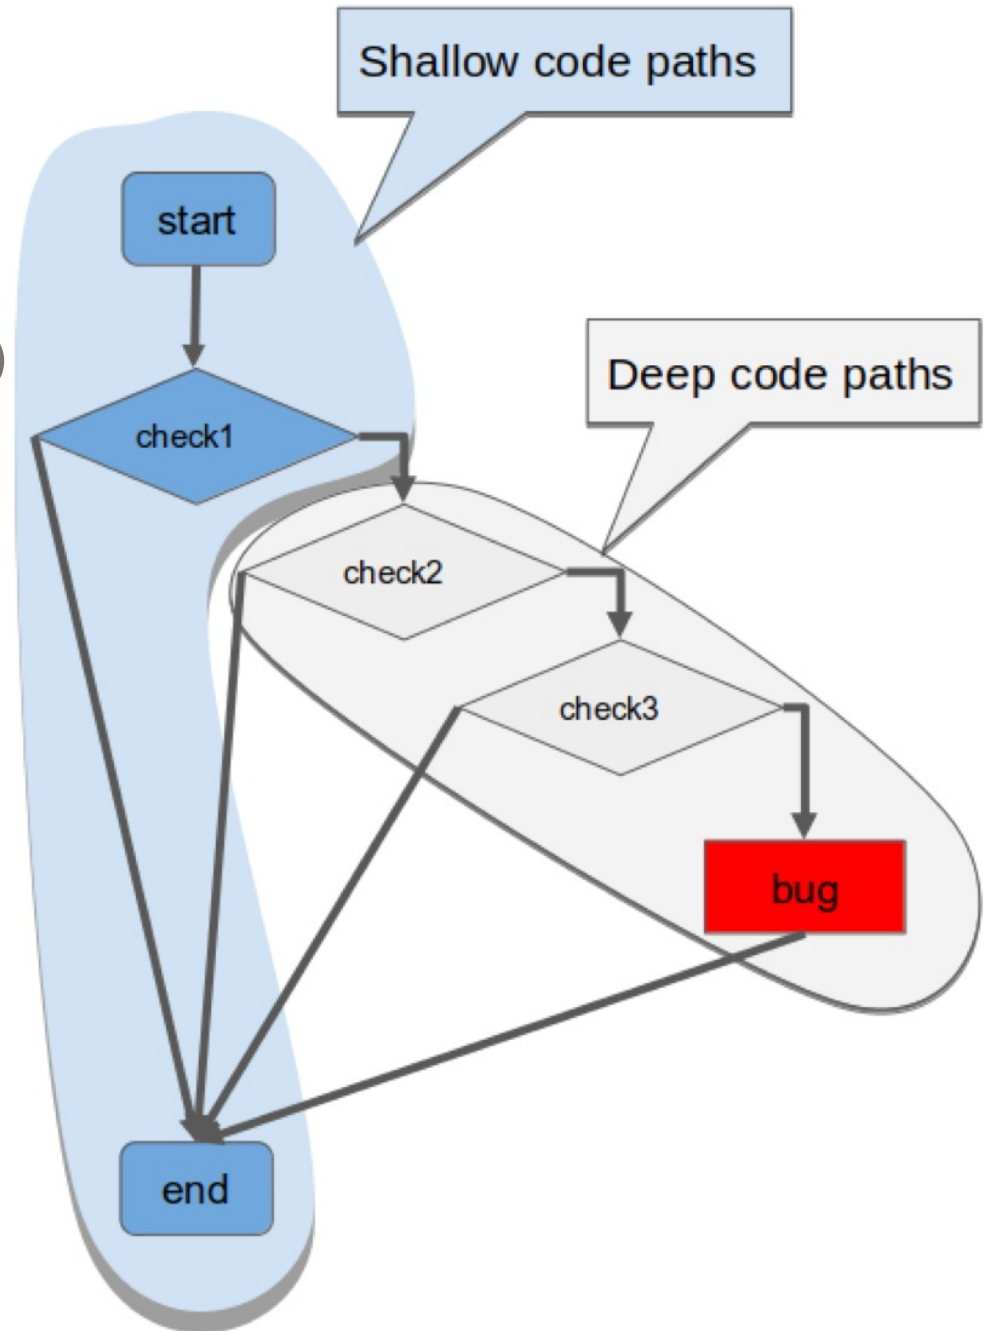
```
int arr[6];
```

```
int func(int a, int b) {  
    int idx = 6;  
    if (a < 6) idx -= 6; else idx -= 1;  
    if (b < 6) idx -= 1; else idx += 1;  
    return arr[idx];  
}
```

- Not all paths are executed: Fail to find overflow bug ($a = 1, b = 1$)
- Full path coverage evaluates all possible paths
 - expensive (path explosion due to each branch)
 - impossible for loops
- Probabilistically covers state space

Coverage wall

- Hard to satisfy checks (e.g., checksum)
- Chains of checks
- Fuzzer no longer makes progress after certain iterations.





How to measure code coverage


Existing tools:

- Gcov:
 - <https://gcc.gnu.org/onlinedocs/gcc/Gcov.html>
- SanitizerCoverage:
 - <https://clang.llvm.org/docs/SourceBasedCodeCoverage.html>



Example

- Compile program with coverage instrumentation
 - `gcc -fprofile-arcs -ftest-coverage gcov_example.c -o example`
- Generate coverage information on the fly
 - `./example`
- Convert coverage information to html report
 - `gcovr --html-detailed coverage.html`



arc: possible branches taken from one basic block to another

```
1  #include <stdio.h>
2  int arr[5] = { 0, 1, 2, 3, 4 };
3
4  int func(int a) {
5      int idx = 5;
6      if (a < 5)
7          idx -= 5;
8      else
9          idx -= 1;
10     return arr[idx];
11 }
12
13 int main(int argc, char *argv[]){
14     if (argc >= 2){
15         printf("argc>=2\n");
16         return 0;
17     }
18     printf("gcov testing.\n");
19     int result = func(argc);
20     return result;
21 }
```




Example

- Generated report:
 - Green line means executed
 - Red line means never executed

GCC Code Coverage Report

Directory: .

File: gcov_example.c

Date: 2022-04-14 17:13:20

	Exec	Total	Coverage
Lines:	10	13	76.9%
Functions:	2	2	100.0%
Branches:	2	4	50.0%

► List of functions

Line	Branch	Exec	Source
1			<code>#include <stdio.h></code>
2			<code>int arr[5] = { 0, 1, 2, 3, 4 };</code>
3			
4		1	<code>int func(int a) {</code>
5		1	<code>int idx = 5;</code>
6	► 1/2	1	<code>if (a < 5)</code>
7		1	<code>idx -= 5;</code>
8			<code>else</code>
9		x	<code>idx -= 1;</code>
10		1	<code>return arr[idx];</code>
11			<code>}</code>
12			
13		1	<code>int main(int argc, char *argv[]){</code>
14	► 1/2	1	<code>if (argc >= 2){</code>
15		x	<code>printf("argc>=2\n");</code>
16		x	<code>return 0;</code>
17			<code>}</code>
18		1	<code>printf("gcov testing.\n");</code>
19		1	<code>int result = func(argc);</code>
20		1	<code>return result;</code>
21			<code>}</code>
22			



Fuzzing

- Feeding in random inputs until the XML parser crashes

— How do we inject inputs? execute the parser with a xml file

— How do we generate inputs? coverage-guided, mutation

— How do we detect bugs?

— How do we automate the process?



Detecting bugs

- Program crash is not a good indicator
 - Some bugs do not cause crash immediately (e.g., memory corruptions)
 - Other bugs are not crashes (e.g., uninitialized value usage)
- Use dynamic bug detectors
 - Sanitizers



Sanitizers

- Tools based on compiler instrumentation.
- Discover bugs like integer overflow, heap buffer overflow, use after free, etc.
- Different type of Sanitizers:
 - ASAN
 - UBSAN
 - MSAN
 - TSAN



Address Sanitizer (ASAN)

- Fast memory error detector
- Compiler directive: `-fsanitize=address`
- Detects various issues:
 - Out-of-bounds access to heap, stack and globals
 - Use After Free
 - Use after scope
- Typical slowdown: $\sim 2x$



Address Sanitizer (ASAN)

```
1  int main(int argc, char **argv) {
2      int *array = new int[100];
3      delete [] array;
4      return array[argc]; // use after free here
5  }
```

compile, link, run

clang++ -g -fsanitize=address example_uaf.cc && ./a.out

```
=====
==38960==ERROR: AddressSanitizer: heap-use-after-free on address 0x00010623a844 at pc 0x000104463f4c bp 0x00016b99f0e0 sp 0x00016b99f0d8
```

```
READ of size 4 at 0x00010623a844 thread T0
```

```
#0 0x104463f48 in main example_uaf.cc:4
```

```
#1 0x1048210f0 in start+0x204 (dyld:arm64+0x50f0)
```

```
#2 0x4d7dfffffffffc (<unknown module>)
```

→ read freed buffer in `return array[argc];`

```
0x00010623a844 is located 4 bytes inside of 400-byte region [0x00010623a840,0x00010623a9d0)
```

```
freed by thread T0 here:
```

```
#0 0x104913c70 in wrap__ZdaPv+0x6c (libcclang_rt.asan_osx_dynamic.dylib:arm64+0x4bc70)
```

```
#1 0x104463efc in main example_uaf.cc:3
```

```
#2 0x1048210f0 in start+0x204 (dyld:arm64+0x50f0)
```

```
#3 0x4d7dfffffffffc (<unknown module>)
```

→ free buffer in `delete [] array;`

```
previously allocated by thread T0 here:
```

```
#0 0x10491387c in wrap__Znam+0x6c (libcclang_rt.asan_osx_dynamic.dylib:arm64+0x4b87c)
```

```
#1 0x104463ee4 in main example_uaf.cc:2
```

```
#2 0x1048210f0 in start+0x204 (dyld:arm64+0x50f0)
```

```
#3 0x4d7dfffffffffc (<unknown module>)
```

→ `int *array = new int[100];`

```
SUMMARY: AddressSanitizer: heap-use-after-free example_uaf.cc:4 in main
```



Undefined Behavior Sanitizer (UBSAN)

- Compiler directive: `-fsanitize=undefined`
- Detects undefined behavior:

- Divide by zero

```
int b = 0;  
int c = a / b;
```

- Signed integer overflow

```
int k = 0x7fffffff;  
k += 1;  
// 0x7fffffff + 1 cannot be represented in type 'int'
```

- Dereferencing misaligned/null pointer

```
char * ptr = (char*)alloc_mem(LARGE_SIZE); // possible NULL pointer  
*ptr = 1;
```



Memory Sanitizer (MSAN)

- Compiler directive: `-fsanitize=memory`
- Detects uninitialized reads

```
1  int main(int argc, char** argv) {  
2      int array[10]; // uninitialized stack array  
3      return array[5];  
4  }
```

`clang -fsanitize=memory msan_example.c`

```
==108578==WARNING: MemorySanitizer: use-of-uninitialized-value  
#0 0x4983bd in main /home/happy/workspace/msan_example.c:3:3  
#1 0x7f93e4ca00b2 in __libc_start_main /build/glibc-sMfBJT/glibc-2.31/csu/../csu/libc-start.c:308:16  
#2 0x41c22d in _start (/home/happy/workspace/a.out+0x41c22d)  
  
SUMMARY: MemorySanitizer: use-of-uninitialized-value /home/happy/workspace/msan_example.c:3:3 in main
```

- Typical slowdown: $\sim 3x$



Thread Sanitizer (TSAN)

- Compiler directive: `-fsanitize=thread`
- Detects data races.

`clang -fsanitize=thread tsan_example.c`

```
1  #include <pthread.h>
2  int g_value;
3  void *thread1(void *x)
4  {
5      g_value = 42; ← write in thread1
6      return x;
7  }
8  int main()
9  {
10     int cnt = 100;
11     while (--cnt > 0)
12     {
13         pthread_t t;
14         pthread_create(&t, NULL, thread1, NULL);
15         g_value = 43; ← write in main thread
16         pthread_join(t, NULL);
17     }
18     return g_value;
19 }
```

```
=====
WARNING: ThreadSanitizer: data race (pid=107809)
  Write of size 4 at 0x000000f18418 by main thread:
    #0 main /home/happy/workspace/tsan_example.c:15:17 (a.out+0x4b5f19)

  Previous write of size 4 at 0x000000f18418 by thread T2:
    #0 thread1 /home/happy/workspace/tsan_example.c:5:13 (a.out+0x4b5e9b)

  Location is global 'g_value' of size 4 at 0x000000f18418 (a.out+0x000000f18418)

  Thread T2 (tid=107812, finished) created by main thread at:
    #0 pthread_create <null> (a.out+0x424b2b)
    #1 main /home/happy/workspace/tsan_example.c:14:9 (a.out+0x4b5f0a)

SUMMARY: ThreadSanitizer: data race /home/happy/workspace/tsan_example.c:15:17 in main
=====
ThreadSanitizer: reported 1 warnings
```

- Typical slowdown: 5x ~ 15x



Fuzzing

- Feeding in random inputs until the XML parser crashes

- How do we inject inputs? Execute the parser with a xml file
- How do we generate inputs? coverage-guided, mutation
- How do we detect bugs? ASAN, UBSAN, MSAN, etc.
- How do we automate the process?



Automation

- Run program with a generated testcase
- Monitor program for coverage and crashes
- Deduplicate crashes
- Minimize testcase
- Generate reproducers
- Report crashes




AFL

- The most well-known fuzzer currently
- Uses compiler-time instrumentation to track branch coverage
- Mutate fuzzing testcase based on previous branch coverage



Core of AFL

- Pseudo code of coverage instrumentation:
 - `cur_location = <COMPILE_TIME_RANDOM>;`
 - `shared_mem[cur_location ^ prev_location] ++;`
 - `prev_location = cur_location >> 1;`
- Different coverage for the following call-chains:
 - `A -> B -> C -> D -> E`
 - `A -> B -> D -> C -> E`








Core of AFL

Corpus Mutation Strategies

- **Bit Flip**: flips a bit. i.e., 1->0, 0->1
 - L/S (length of toggled bits/stepover bits): 1/1, 2/1, 4/1, 8/8...
- **Byte Flip**: flips a byte
- **Arithmetic**: subtract/add small integer to 8/16/32 bit values.
- **Havoc**: random things with bit/byte/arithmetic, etc.
- **Interest**: replace content with **known interesting value** (e.g., 65535)
- **Dictionary**: user provided dictionary or auto discovered tokens.
- **Splice**: split & combine two or more files to get a new file



Hands On: Building AFL

- git clone <https://github.com/google/AFL.git>
 - make
 - cd llvm_mode
 - make
- | | | |
|---|--|---|
| { |  afl-clang-fast.c | wrapper for clang |
| |  afl-llvm-pass.so.cc | Instrumentation pass for clang,
add coverage support for program |
| |  afl-llvm-rt.o.c | Implementation of instrument function |
- **afl-clang-fast**: compiler which instruments program while building
 - **afl-fuzz**: overall fuzzer



Hands On: Fuzzing example

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  #define SIZE 20
5
6  int main(int argc, char *argv[])
7  {
8      char input[SIZE] = {0};
9      size_t length;
10     length = read(STDIN_FILENO, input, SIZE);
11     char calc_value = 0;
12
13     if (input[0] == 'a'){
14         if (input[1] == 'b'){
15             if (input[2] == 'c'){
16                 int idx = input[3] + input[4];
17                 calc_value = input[idx];
18             }
19         }
20     }
21
22     return calc_value;
23 }
```

buffer overflow bug

- Example PoC:

input buffer: "abc\x40\x40"



Hands On: Fuzzing example

- Build program with **ASan & Coverage** instrumentation
 - afl-clang-fast -fsanitize=address example.c

```
21 if ( !_asan_option_detect_stack_use_after_return )
22 {
23     v4 = 0LL;
24     goto LABEL_3;
25 }
26 _asan_stack_malloc_1();
27 v4 = v3;
28 v5 = (unsigned __int64)v3;
29 if ( !v3 )
30 LABEL_3:
31     v5 = (unsigned __int64)&v19[-12];
32     v19[3] = v5;
33     *(_QWORD *)v5 = 1102416563LL;
34     *(_QWORD *)(v5 + 8) = "1 32 20 7 input:8";
35     *(_QWORD *)(v5 + 16) = main;
36     v6 = v5 >> 3;
37     *(_QWORD *)(v6 + 2147450880) = 0xF3F8F8F8F1F1F1F1LL;
38     *(_DWORD *)(v6 + 2147450888) = -202116109;
39     v7 = (_BYTE *) (v5 + 32);
40     _afl_area_ptr[_start__sancov_guards] += __CFADD__( _afl_area_ptr[_start__sancov_guards], 1) + 1;
41     *(_WORD *) (v6 + 2147450884) = 0;
42     *(_BYTE *) (v6 + 2147450886) = 4;
43     _asan_memset();
44     v8 = (_QWORD *) (v5 + 32);
45     _interceptor_read();
46     v9 = *(_BYTE *) (((v5 + 32) >> 3) + 0x7FFF8000);
47     if ( v9 && ((unsigned __int8)v7 & 7) >= v9 )
48     {
49         v10 = v5 + 32;
50         v11 = _asan_report_load1(v5 + 32, v8);
51     }
52     else
53     {
54         if ( *v7 != 97 )
55         {
```

} Instead of alloc array on stack,
ASan wrap the alloc on FakeStack

← Increase coverage,
add counter to shared memory



Hands On: Fuzzing example

- Build program with **ASan & Coverage** instrumentation
 - afl-clang-fast -fsanitize=address example.c

```
21 if ( !_asan_option_detect_stack_use_after_return )
22 {
23     v4 = 0LL;
24     goto LABEL_3;
25 }
```

```
26 _asan_stack_malloc_1();
27 v4 = v3;
28 v5 = (unsigned __int64)v3;
29 if ( !v3 )
```

} Instead of alloc array on stack,
ASan wrap the alloc on FakeStack

```
30 LABEL_3:
31     v5 = (unsigned __int64)&v19[-12];
32     v19[3] = v5;
33     *(_QWORD *)v5 = 1102416563LL;
34     *(_QWORD *)(v5 + 8) = "1 32 20 7 input:8";
35     *(_QWORD *)(v5 + 16) = main;
36     v6 = v5 >> 3;
37     *(_QWORD *)(v6 + 2147450880) = 0xF3F8F8F8F1F1F11LL;
38     *(_DWORD *)(v6 + 2147450888) = -202116109;
39     v7 = (_BYTE *)(v5 + 32);
```

```
40     _afl_area_ptr[_start__sancov_guards] += __CFADD__(_afl_area_ptr[_start__sancov_guards], 1) + 1;
41     *(_WORD *)(v6 + 2147450884) = 0;
42     *(_BYTE *)(v6 + 2147450886) = 4;
```

```
43     _asan_memset();
44     v8 = (_QWORD *)(v5 + 32);
45     _interceptor_read();
46     v9 = *(_BYTE *)(((v5 + 32) >> 3) + 0x7FFF8000);
47     if ( v9 && ((unsigned __int8)v7 & 7) >= v9 )
48     {
49         v10 = v5 + 32;
50         v11 = _asan_report_load1(v5 + 32, v8);
51     }
```

```
52 else
53 {
54     if ( *v7 != 97 )
55     {
```

← Increase coverage,
add counter to shared memory

← Check buffer sanity (i.e., input[0]) before read



Hands On: Fuzzing example

- Build program with **ASan** & **Coverage** instrumentation
 - `afl-clang-fast -fsanitize=address example.c`
- Provide initial corpus/testcase seed
 - `echo 123 > corpus/seeds0`
- Fuzz it
 - `afl-fuzz -D -i input -o output -- ./a.out`
 - `-D`: enable deterministic fuzzing (more mutation strategies)
 - `-i dir`: input directory with initial testcases.
 - `-o dir`: output directory for fuzzer findings



Hands On: Fuzzing example

One cycle:

- go over all the interesting test cases discovered so far

```
american fuzzy lop 2.57b (a.out)

process timing
  run time : 0 days, 0 hrs, 0 min, 3 sec
  last new path : 0 days, 0 hrs, 0 min, 1 sec
  last uniq crash : 0 days, 0 hrs, 0 min, 2 sec
  last uniq hang : none seen yet

cycle progress
  now processing : 3 (75.00%)
  paths timed out : 0 (0.00%)

stage progress
  now trying : havoc
  stage execs : 4888/12.3k (39.78%)
  total execs : 19.7k
  exec speed : 5237/sec

fuzzing strategy yields
  bit flips : 0/120, 1/116, 1/108
  byte flips : 0/15, 0/11, 0/3
  arithmetics : 0/840, 0/160, 0/0
  known ints : 0/70, 0/298, 0/132
  dictionary : 0/0, 0/0, 0/0
  havoc : 2/9152, 0/3792
  trim : 50.00%/1, 0.00%

map coverage
  map density : 0.01% / 0.01%
  count coverage : 1.00 bits/tuple

findings in depth
  favored paths : 1 (25.00%)
  new edges on : 4 (100.00%)
  total crashes : 2 (1 unique)
  total tmouts : 0 (0 unique)

path geometry
  levels : 4
  pending : 1
  pend fav : 1
  own finds : 3
  imported : n/a
  stability : 100.00%

overall results
  cycles done : 5
  total paths : 4
  uniq crashes : 1
  uniq hangs : 0

[cpu000: 16%]
```

branch tuples already hit/the bitmap can hold

the number of crashes

consistency of observed traces



Hands On: Fuzzing example

Output (fuzzer status & findings):

output

```
├── crashes
│   ├── id:000000,sig:06,src:000002,op:flip4,pos:2
│   └── README.txt
├── fuzz_bitmap
├── fuzzer_stats
├── hangs
├── plot_data
└── queue
    ├── id:000000,orig:seeds0
    ├── id:000001,src:000000,op:havoc,rep:64,+cov
    ├── id:000002,src:000001,op:flip2,pos:1,+cov
    └── id:000003,src:000002,op:havoc,rep:4,+cov
```

← testcase which crashes the program

← fuzzing status. For clustered fuzzers, use **afl-whatsup** to get status of fuzzers

} interesting testcases discovered

3 directories, 9 files



Hands On: Fuzzing example

Reproduce crashes:

```
% hexdump -C output/crashes/id:000000,sig:06,src:000002,op:flip4,pos:2
00000000  61 62 63 ff                                |abc.| ← "abc\xff"
```

```
$ cat output/crashes/id:000000,sig:06,src:000002,op:flip4,pos:2| ./a.out
=====
==4016427==ERROR: AddressSanitizer: stack-buffer-underflow on address 0x7ffee926c62f at pc 0x555c95a4e729 bp 0x7ffee926c600 sp 0x7ffee926c5f0
READ of size 1 at 0x7ffee926c62f thread T0
#0 0x555c95a4e728 in main /home/happy/fuzz/afl_lab/example.c:17
#1 0x7ff1904090b2 in __libc_start_main (/lib/x86_64-linux-gnu/libc.so.6+0x240b2)
#2 0x555c95a4e7ad in _start (/home/happy/fuzz/afl_lab/a.out+0x17ad)

Address 0x7ffee926c62f is located in stack of thread T0 at offset 31 in frame
#0 0x555c95a4e22f in main /home/happy/fuzz/afl_lab/example.c:7

This frame has 1 object(s):
[32, 52) 'input' (line 8) ← Memory access at offset 31 underflows this variable
HINT: this may be a false positive if your program uses some custom stack unwind mechanism, swapcontext or vfork
(longjmp and C++ exceptions *are* supported)
```



Other fuzzers

- **libFuzzer**
 - In-process, coverage-guided fuzzing engine.
- **LibAFL**
 - Advanced fuzzing library written in Rust
 - Scales across cores and machines: Windows, Android, MacOS, Linux, no_std, etc.
- **AFL++**
 - superior fork to Google's AFL
 - more speed, more and better mutations, instrumentation, custom module support, etc.
- **Nautilus**
 - A grammar based feedback Fuzzer



Summary

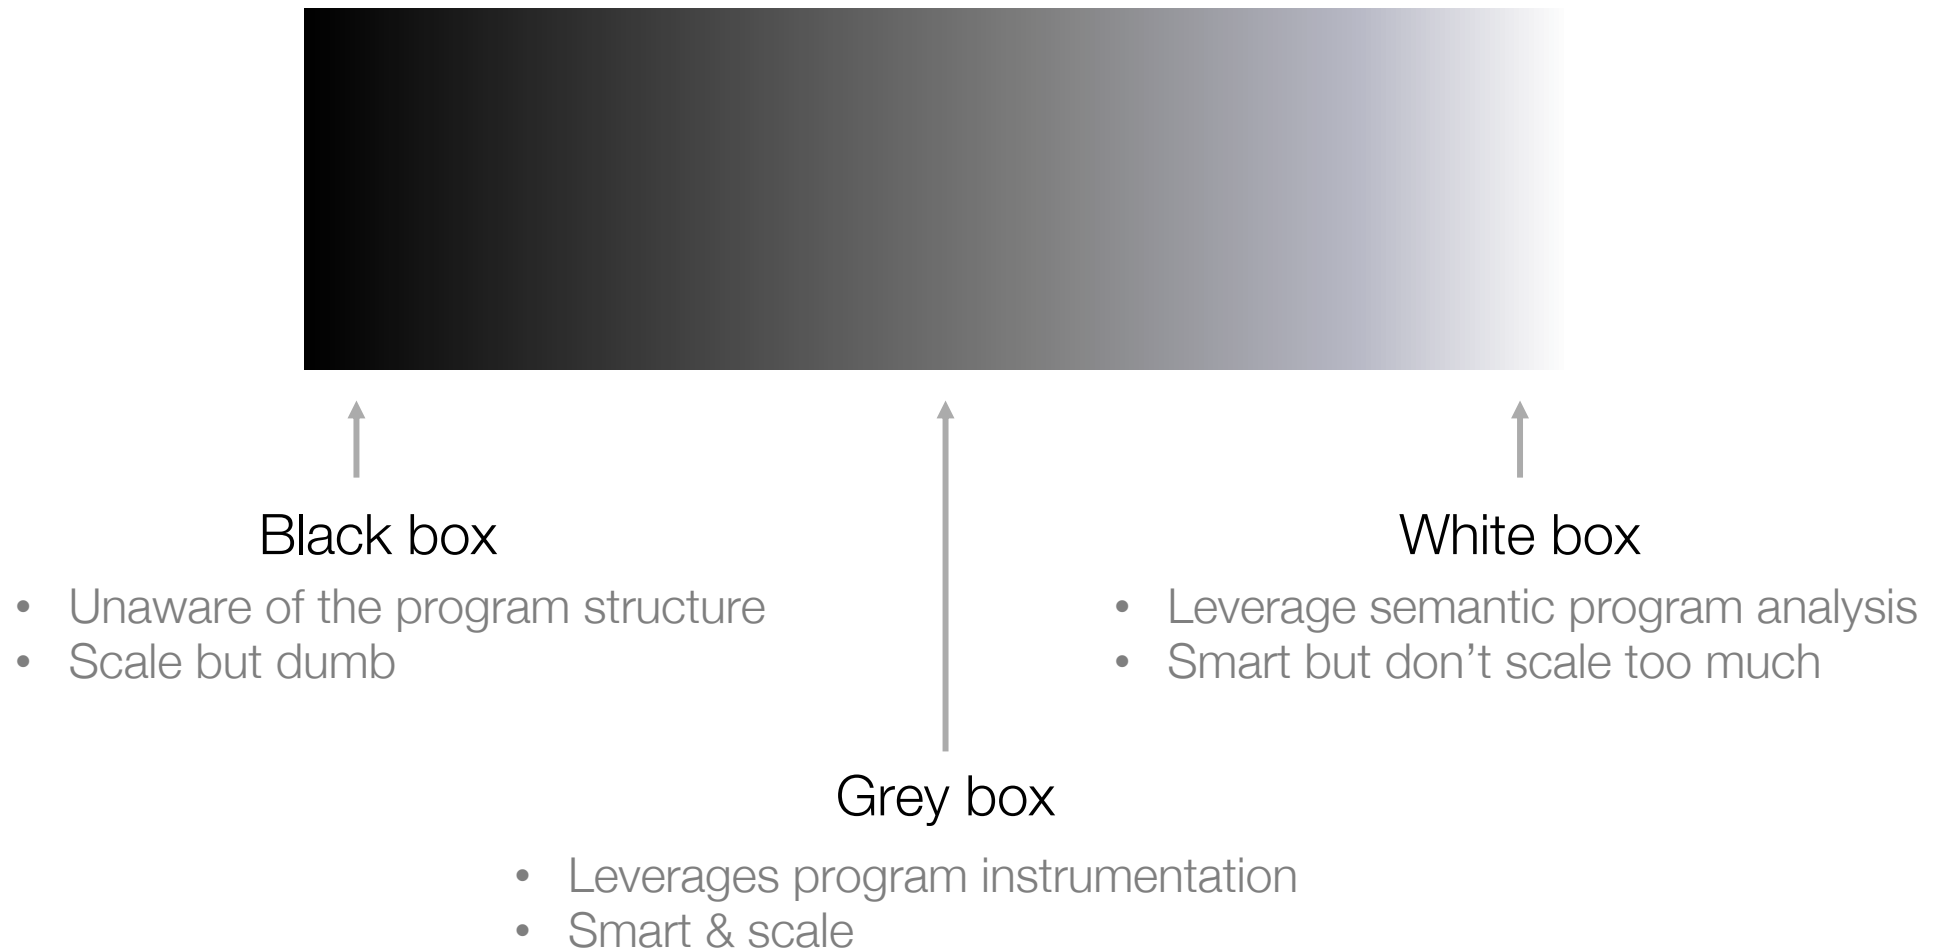
- Feeding in random inputs until the program crashes

- How do we inject inputs? Execute the parser with a xml file
- How do we generate inputs? coverage-guided, mutation
- How do we detect bugs? ASAN, UBSAN, MSAN, etc.
- How do we automate the process? AFL, libFuzzer, etc.



Conclusion

Three Shades of Fuzzing





Conclusion

Principle of Greybox Fuzzing

1. Preprocess
2. Scheduling: Choose "good" inputs
3. Input Generation: Mutations
4. Input Evaluation: Observe coverage/score
5. Configuration Updating
6. Continue