



Control-Flow Integrity

Yajin Zhou (<http://yajin.org>)

Zhejiang University



Motivation

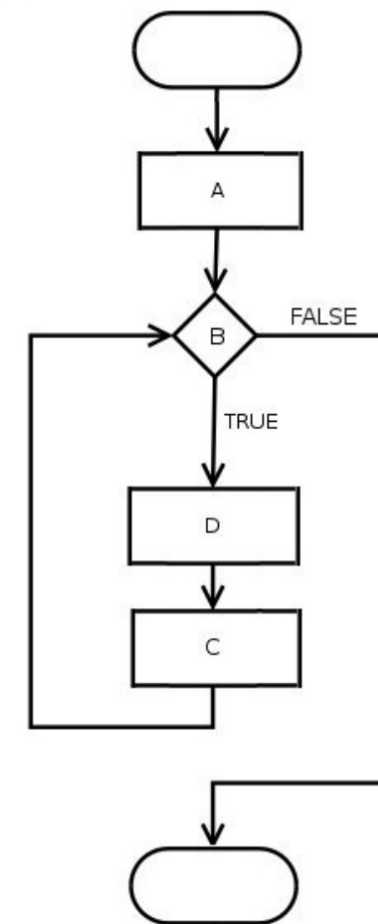
- Code injection
 - W^X
- Code reuse
 - ASLR
 - CFI

Program control flow

- Unconditional jumps
- Conditional jumps
- Loops
- Subroutines
- Unconditional halt

for(A;B;C)

D;





vuln.c

```
#include <stdio.h>
#include <string.h>

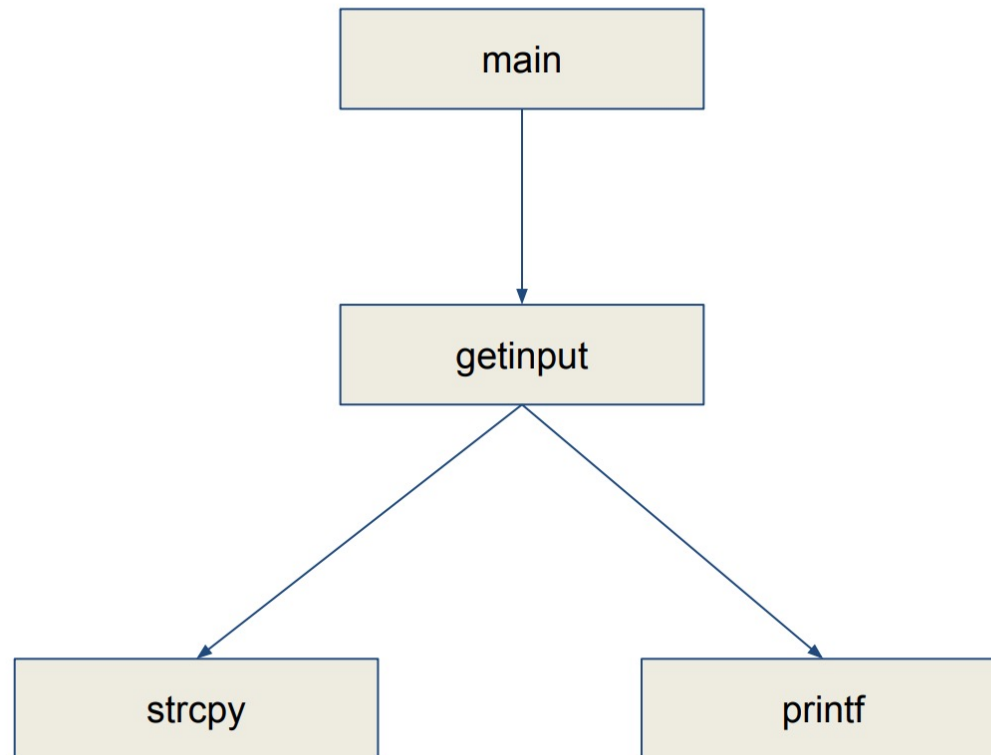
void getinput(char *input) {
    char buffer[32];

    strcpy(buffer, input);
    printf("You entered: %s\n", buffer);
}

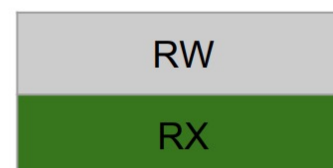
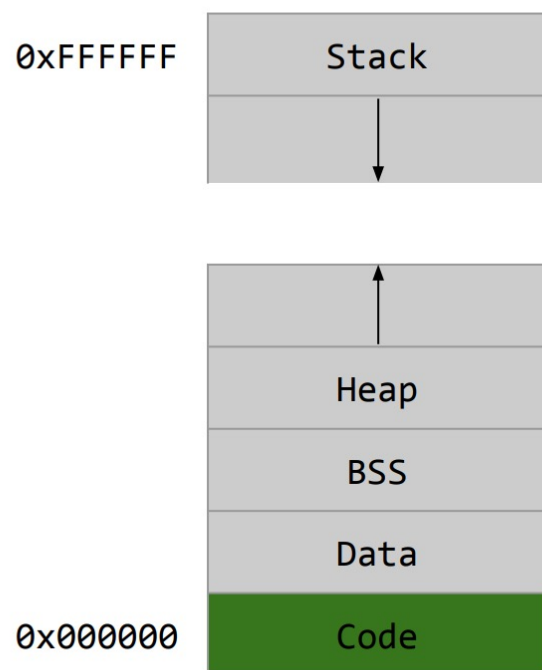
int main(int argc, char **argv) {
    getinput(argv[1]);
    return(0);
}
```



Simple call graph

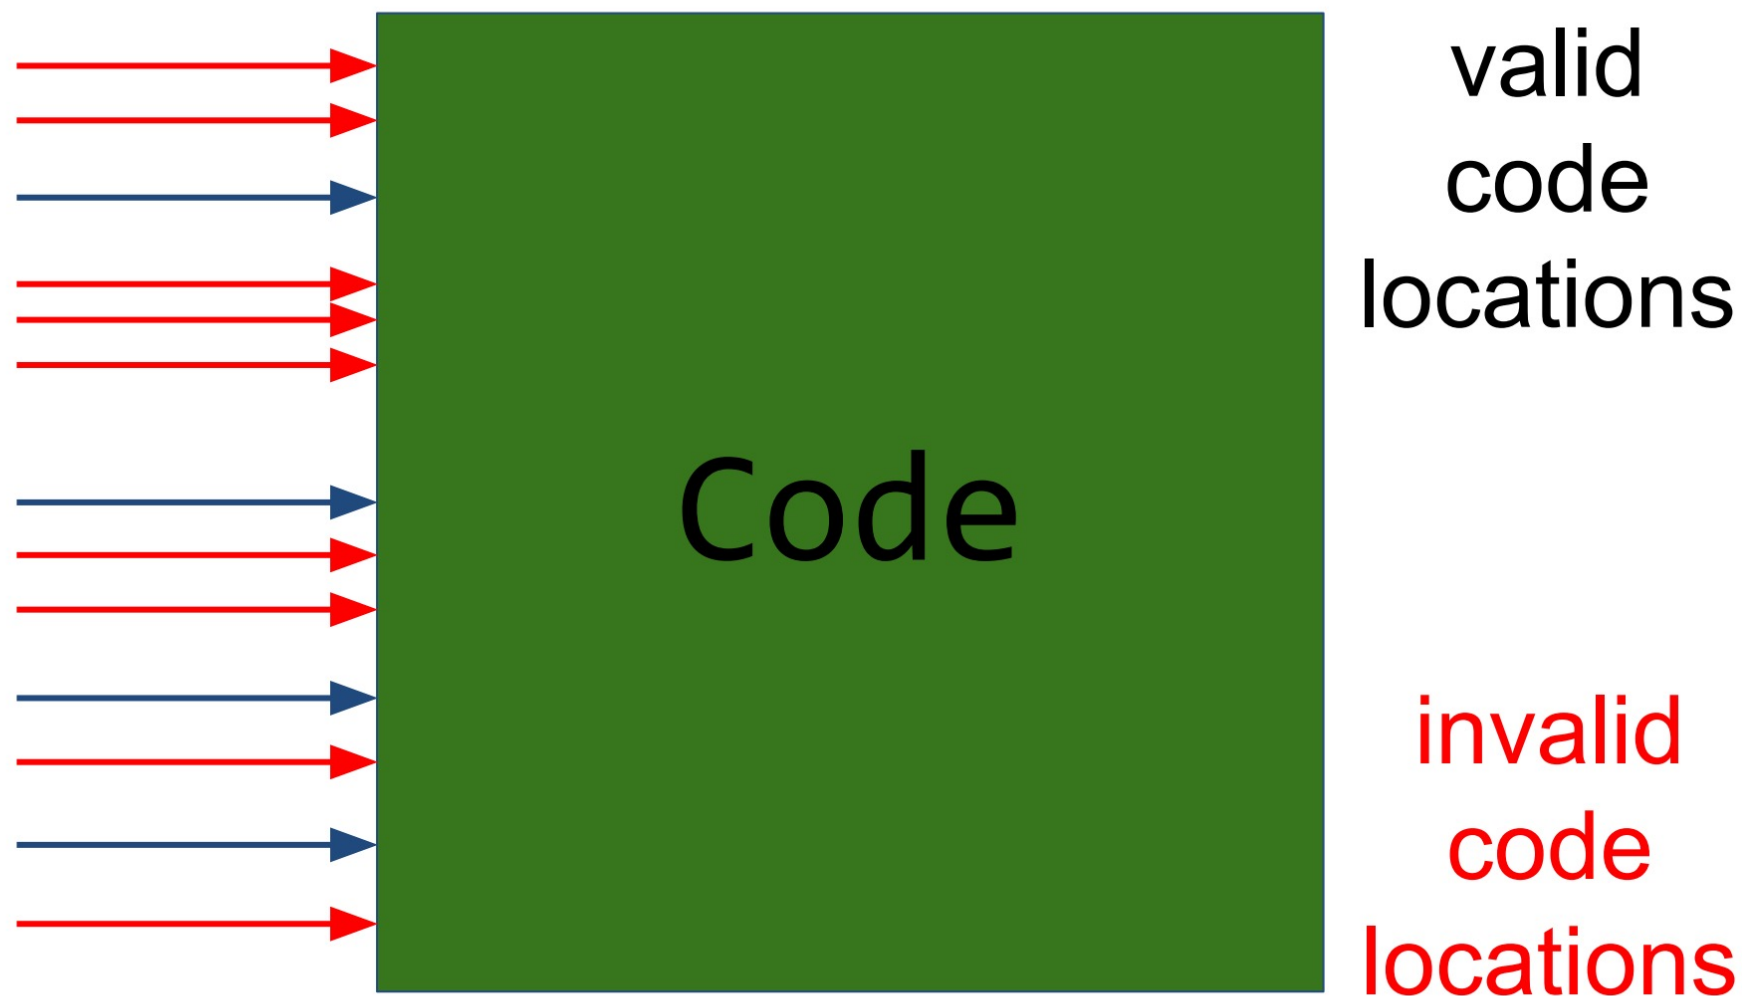


W^X





ROP





ROP's control stack (just data)

The diagram illustrates the ROP's control stack, which is composed of just data. It shows a sequence of assembly instructions, with specific instructions highlighted in red boxes and arrows indicating the control flow between them.

```

movl $0x0000002e,%eax,%eax
movl $0x004(%ebp),%eax
movl %eax,%eax
call 0x008b11
add $0x1f,%eax
and $0x1f,%eax
subl %eax,%esp
ssll $0x2(%esp),%edx
movl %edx,%eax,%esp
jmp 0x0006b3b4
incl 0x4(%ebp)
movl 0x4(%ebp),%eax
movzbl %eax,%ecx
cmpl $0x3a,%cl
je 0x0006b3b1
movl 0xb4(%ebp),%ebx
jne 0x0006b3bb
movl $0x00,%eax
movl $0x0,%edi(%ebp)
jmp 0x0006b3bd
movb %edi,%eax
incl %eax
incl 0x4(%ebp)
movl 0x4(%ebp),%eax
movzbl %eax,%ecx
testb %ecx,%cl
setnl %cl
cmpl $0x3a,%cl
je 0x0006b3d1
testb %al,%cl
jne 0x0006b3d1
movl 0x0(%esp),%eax
cmpl $0x01,0x0008a780
jne 0x0006b3d1
movl 0x4(%esp),%eax
movl $0x00000021,0x4(%esp)
movl %eax,%esp
call 0x0006b3b4
testl %eax,%eax
jre 0x0006b3b4
movl 0xb4(%ebp),%esi
movl $0x00000002,%ecx
movl $0x00007e270,%edi
cdd
repz/repnb (%esi),(%edi)
movl $0x00000000,%eax
je 0x0006b3de
movzbl 0xf(%esi),%eax
movzbl 0xf(%edi),%eax
subl %ecx,%eax
testl %eax,%eax
jrl 0x0006b3e3
movl 0xb4(%ebp),%esi
movl $0x0000700bb,%eax
movl $0x00000000,%ecx
repz/repnb (%esi),(%ecx)
je 0x0006b3e6
movzbl 0xf(%esi),%ecx
movzbl 0xf(%edi),%ecx
subl %ecx,%ecx
testl %ecx,%ecx

```

Arrows indicate the control flow between the highlighted instructions, showing a sequence of jumps and conditional jumps that form the ROP chain.

Code is not executed in the intended way!

Control-Flow Integrity: CCS 05



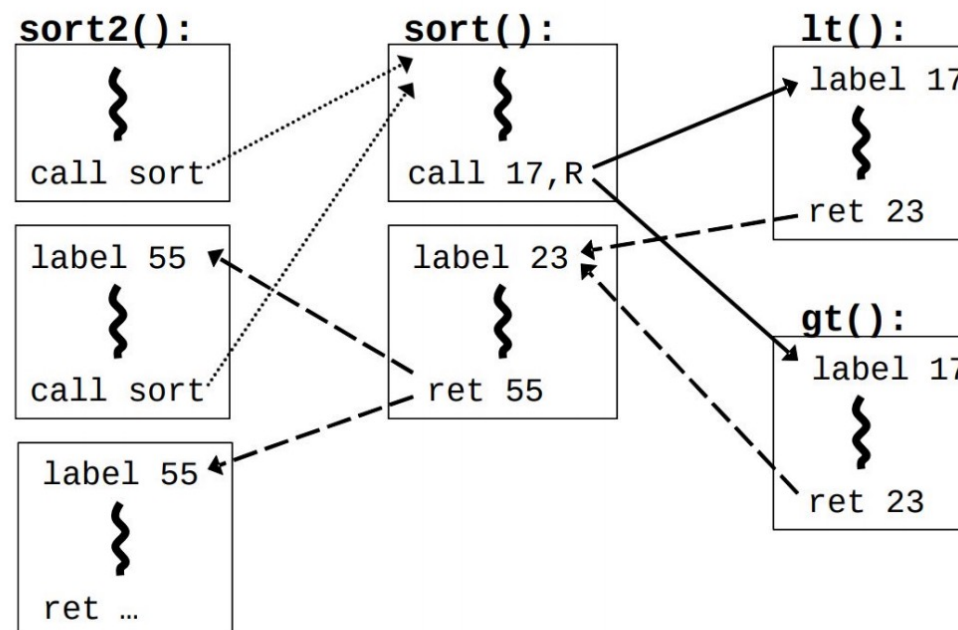
Control-flow integrity

- CFI is a security policy
- Execution must follow a path of a Control-Flow Graph
- CFG can be pre-computed
 - source-code analysis
 - binary analysis
 - execution profiling
- Forward-edge and backward-edge
- But how can we enforce this extracted control-flow?



Enforcing CFI by Instrumentation

```
bool lt(int x, int y) {  
    return x < y;  
}  
  
bool gt(int x, int y) {  
    return x > y;  
}  
  
sort2(int a[], int b[], int len)  
{  
    sort( a, len, lt );  
    sort( b, len, gt );  
}
```



- LABEL ID
- CALL ID, DST
- RET ID



CFI Instrumentation Code

Source		Destination	
Opcode bytes	Instructions	Opcode bytes	Instructions
FF E1	jmp ecx ; computed jump	8B 44 24 04	mov eax, [esp+4] ; dst
		...	

can be instrumented as (a):

81 39 78 56 34 12	cmp [ecx], 12345678h ; comp ID & dst	78 56 34 12	; data 12345678h ; ID
75 13	jne error_label ; if != fail	8B 44 24 04	mov eax, [esp+4] ; dst
8D 49 04	lea ecx, [ecx+4] ; skip ID at dst	...	
FF E1	jmp ecx ; jump to dst		

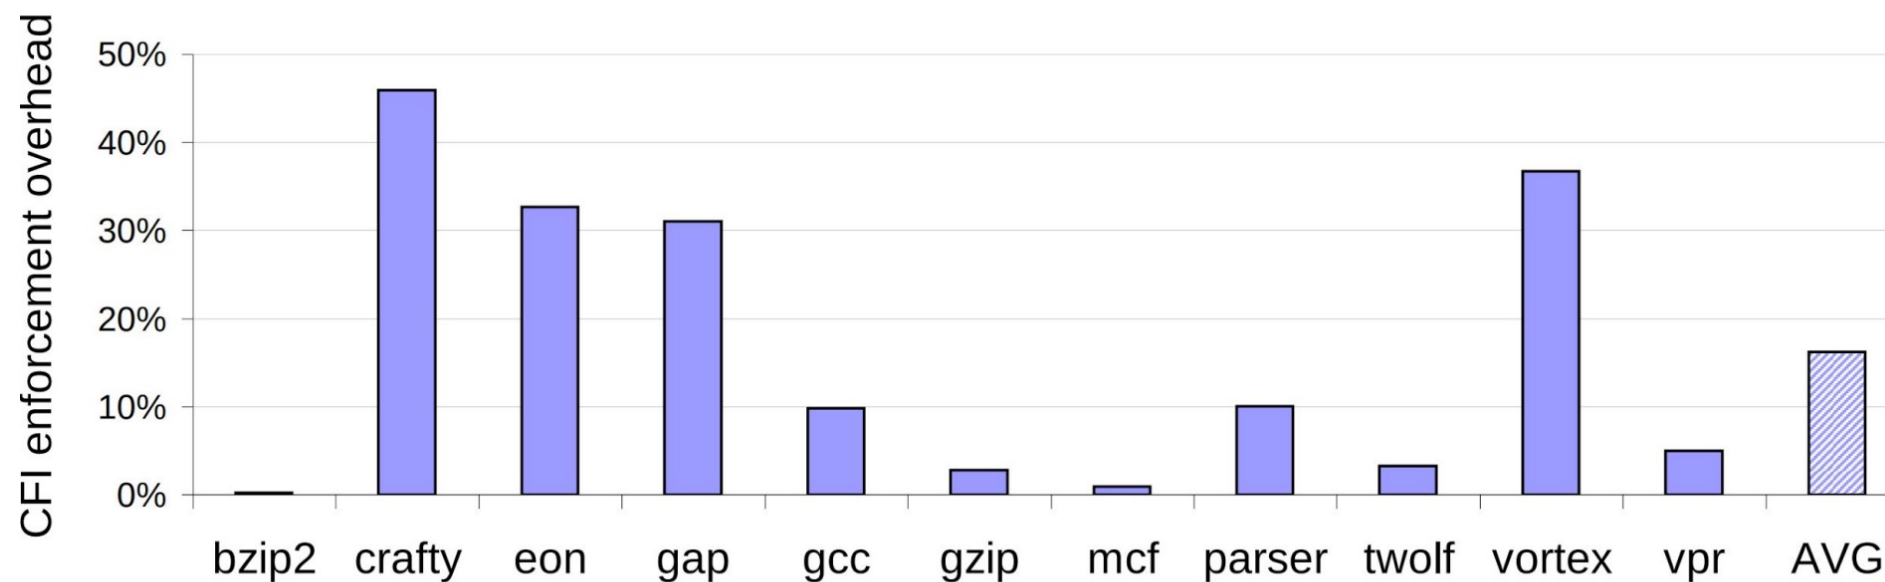
or, alternatively, instrumented as (b):

B8 77 56 34 12	mov eax, 12345677h ; load ID-1	3E 0F 18 05	prefetchnta ; label
40	inc eax ; add 1 for ID	78 56 34 12	[12345678h] ; ID
39 41 04	cmp [ecx+4], eax ; compare w/dst	8B 44 24 04	mov eax, [esp+4] ; dst
75 13	jne error_label ; if != fail	...	
FF E1	jmp ecx ; jump to label		

- The extra code checks that the destination code is the intended jump location



Overhead





Limitation

- Overhead is high
- Precise CFG construction is hard (or even impossible)



CFI in real systems

- Control Flow Guard
- LLVM
- Hardware features
 - Shadow stack
 - IBT: indirect branch tracking



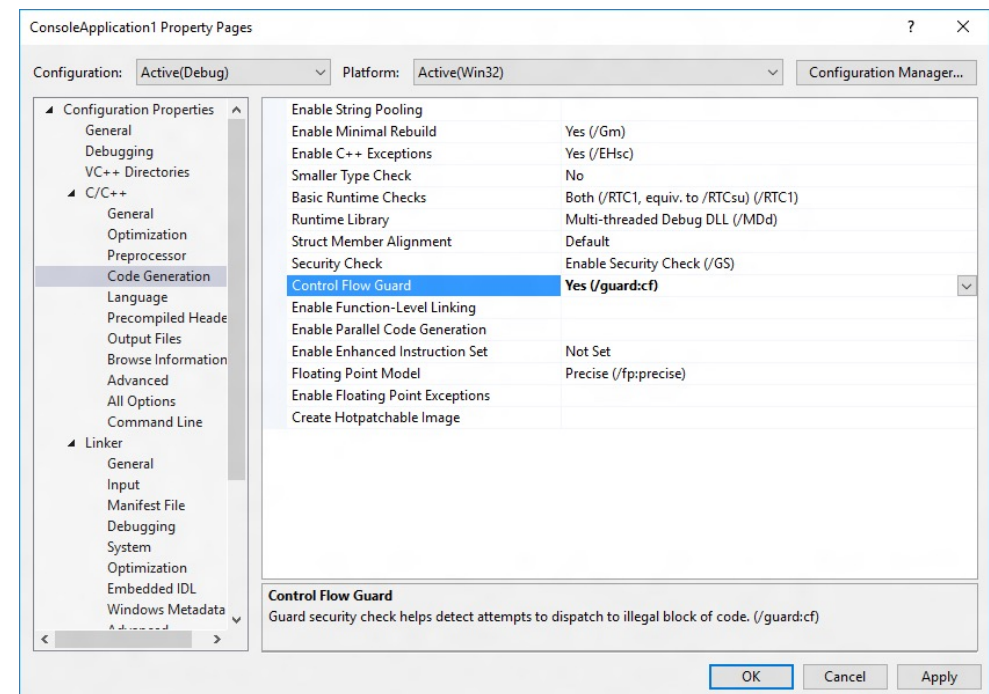
Control Flow Guard

- Windows 10 and Windows 8.1
- Microsoft Visual Studio 2015+
- Adds lightweight security checks to the compiled code
- Identifies the set of functions in the application that are valid targets for indirect calls
- The runtime support, provided by the Windows kernel:
 - Efficiently maintains state that identifies valid indirect call targets
 - Implements the logic that verifies that an indirect call target is valid



Control Flow Guard

- In most cases, there is no need to change source code. All you have to do is add an option to your Visual Studio 2015 project, and the compiler and linker will enable CFG.
- The simplest method is to navigate to **Project | Properties | Configuration Properties | C/C++ | Code Generation** and choose **Yes (/guard:cf)** for Control Flow Guard.



source: Control Flow Guard ([link](#))

How Do I Tell That a Binary is under Control Flow Guard?

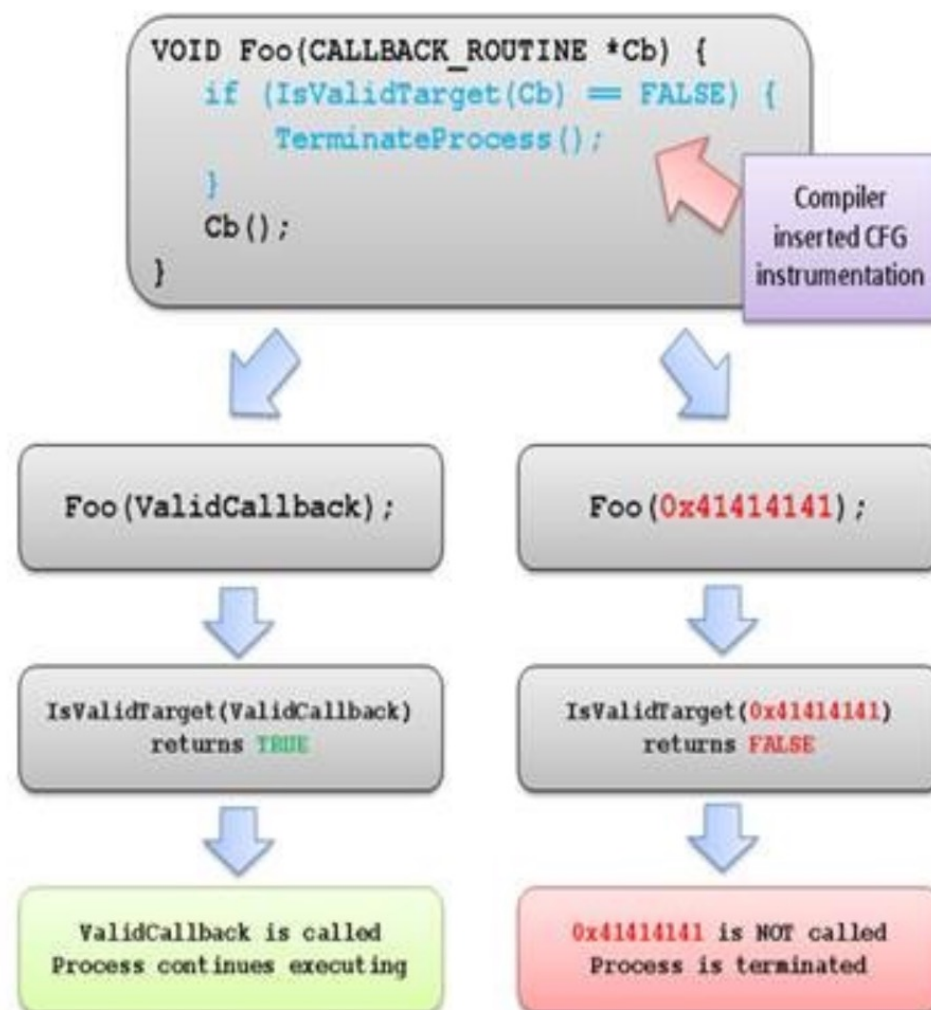


- Run the [dumpbin tool](#) (included in the Visual Studio 2015 installation) from the Visual Studio command prompt with the */headers* and */loadconfig* options: `dumpbin /headers`.
- "CF Instrumented" and "FID table present".

```
Section contains the following load config:

000000A0 size
0 time date stamp
0.00 Version
0 GlobalFlags Clear
0 GlobalFlags Set
0 Critical Section Default Timeout
0 Decommit Free Block Threshold
0 Decommit Total Free Threshold
0000000000000000 Lock Prefix Table
0 Maximum Allocation Size
0 Virtual Memory Threshold
0 Process Heap Flags
0 Process Affinity Mask
0 CSD Version
0000 Reserved
0000000000000000 Edit list
000000014023C008 Security Cookie
00000001401C41A0 Guard CF address of check-function pointer
00000001401C41A8 Guard CF address of dispatch-function pointer
00000001401C42A8 Guard CF function table
E95 Guard CF function count
00003500 Guard Flags
CF Instrumented
FID table present
Protect delayload IAT
Delayload IAT in its own section
```

source: Control Flow Guard ([link](#))





LLVM-CFI

- Clang includes an implementation of a number of control flow integrity (CFI) schemes
- Relies on link-time optimization: LTO
- LLVM-CFI implements the CFI check as a range check that maps into a table of trampolines. All **address taken functions of the same type are translated into a range where they are placed next to each other**. A CFI dispatch is then matched into a jump into an 8-byte aligned jump into the area of targets for the corresponding type.

```
void (*func)();

// func either points to bar or baz
if (usr == MAGIC)
    func = bar;
else
    func = baz;
```



```
void bar();
void baz();
void buz();
void bez(int, int);

void foo(int usr) {
    void (*func)();

    // func either points to bar or baz
    if (usr == MAGIC)
        func = bar;
    else
        func = baz;

    // forward edge CFI check
    // depending on the precision of CFI:
    // a) all functions {bar, baz, buz, bez, foo} are allowed
    // b) all functions with prototype "void (*)()" are allowed,
    //    i.e., {bar, baz, buz}
    // c) only address taken functions are allowed, i.e., {bar, baz}
    CHECK_CFI_FORWARD(func);
    func();

    // backward edge CFI check
    CHECK_CFI_BACKWARD();
}
```



Indirect branch tracking

- ENDBRANCH -> new CPU instruction
- marks valid indirect call/jmp targets in the program
- the CPU implements a state machine that tracks indirect jmp and call instructions
- when one of these instructions is seen, the state machine moves from IDLE to WAIT_FOR_ENDBRANCH state
- if an ENDBRANCH is not seen the processor causes a control protection fault



Challenges

- CFG construction
- A CFG is a graph that covers all valid executions of the program. Nodes in the graph are locations of control-flow transfers in the program and edges encode reachable targets.
- For forward edges, the CFG generation enumerates all possible targets, often leveraging information from the underlying source language.



Indirect functions

- For indirect function calls through a function pointer, the underlying analysis becomes more complicated as the target may not be known a-priori.
- Common source-based analyses use a type-based approach and, looking at the function prototype of the function pointer that is used, enumerate all matching functions.
 - Use address-taken to reduce valid set
 - For virtual calls, i.e., indirect calls in C++ that depend on the type of the object and the class relationship, the analysis can further leverage the type of the object to restrict the valid functions



Indirect functions

- So far, the constructed CFG is **stateless**, i.e., the *context of the execution* is not considered and each control-flow transfer is independent of all others.
- This over-approximates the number of valid targets with different granularities – since at runtime only one target is allowed for any possible transfer



Adaptive Call-site Sensitive Control Flow Integrity

```
1  typedef int (*Handler)(char *);
2  int proceed(Handler handler, char *root_path)
3  {
4      ...
5      return handler(root_path);
6  }
7
8  void auth()
9  {
10     char *user_name;
11     char *passwd;
12     char id[80];
13     int attempt = 5;
14     Handler handler;
15
16     while (attempt > 0) {
17         handler = &on_failure;
18         username = passwd = null;
19
20         scanf("%ms", &username);
21         scanf("%ms", &password);
22
23         passwd = salt_passwd(username, passwd);
24         sprintf(id, "%s;%s", username, passwd);
25
26         if (is_admin(id)) {
27             handler = &on_admin;
28             proceed(handler, user_home_dir);
29         } else {
30             proceed(handler, "/tmp");
31         }
32         attempt--;
33
34         //clear passwd, free user_name, passwd
35         ...
36     }
37 }
```

CFI-LB

- CFI-LB enforces a call-site sensitive CFI policy in which the targets of an indirect branch are validated in the context of callsites (i.e., return addresses on the stack).

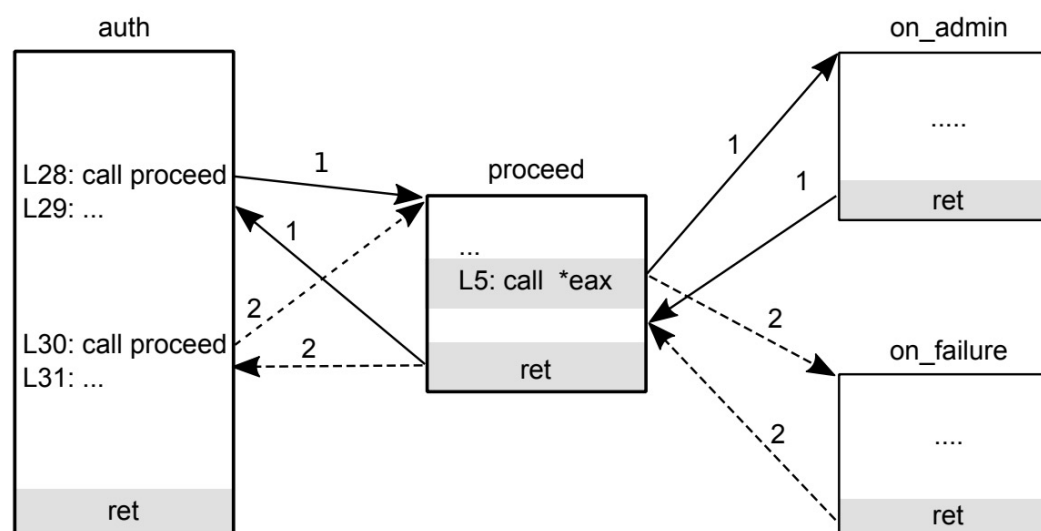


Fig. 3. CFG for the example in Fig. 2



Discussion

- Shared library
- Binary only
- Hardware assisted